

Data Refinery Lab

This lab will use the Titanic data set to demonstrate data profiling, data visualization, and data preparation capabilities of the Data Refinery tool. The lab consists of the following steps:

1. Use the Data Refinery Tool to:
 - a. Profile the data to help determine missing values
 - b. Visualize the data to gain a better understanding
 - c. Prepare the data for modeling
 - d. Run the sequence of data preparation operations on the entire data set.

Step 1: Adding a Data Asset to the DSX Local Labs project.

1. Download a file that contains a random sample of the Titanic data file (note the Data Refinery currently operates on 100 records of a dataset for interactive operation) by clicking on the link [Titanic Random Data Set](#) and following the instructions below.

102 lines (101 sloc) | 10.4 KB

Raw


Blame


History

Search this file...

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare
2	1	0	Allison, Miss. Helen Loraine	female	2.000000	1	2	113781	151.550000
3	1	1	Andrews, Miss. Kornelia Theodosia	female	63.000000	1	0	13502	77.958300
4	1	1	Barkworth, Mr. Algernon Henry Wilson	male	80.000000	0	0	27042	30.000000
5	1	1	Bonnell, Miss. Caroline	female	30.000000	0	0	36928	164.866700
6	1	1	Carter, Miss. Lucile Polk	female	14.000000	1	2	113760	120.000000
7	1	1	Duff Gordon, Sir. Cosmo Edmund (Mr Morgan)	male	49.000000	1	0	PC 17485	56.929200
8	1	1	Fortune, Mrs. Mark (Mary McDougald)	female	60.000000	1	4	19950	263.000000

Right click on Raw and click on Save link as


2. Go to the DSX Local project. Click on the hamburger icon , and then click on Projects and then the DSX_Demos_XX project.



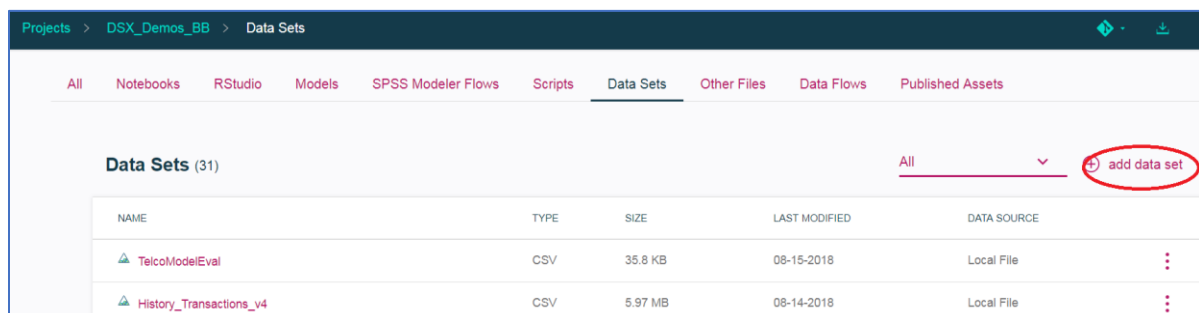
3. Click on **Assets**.



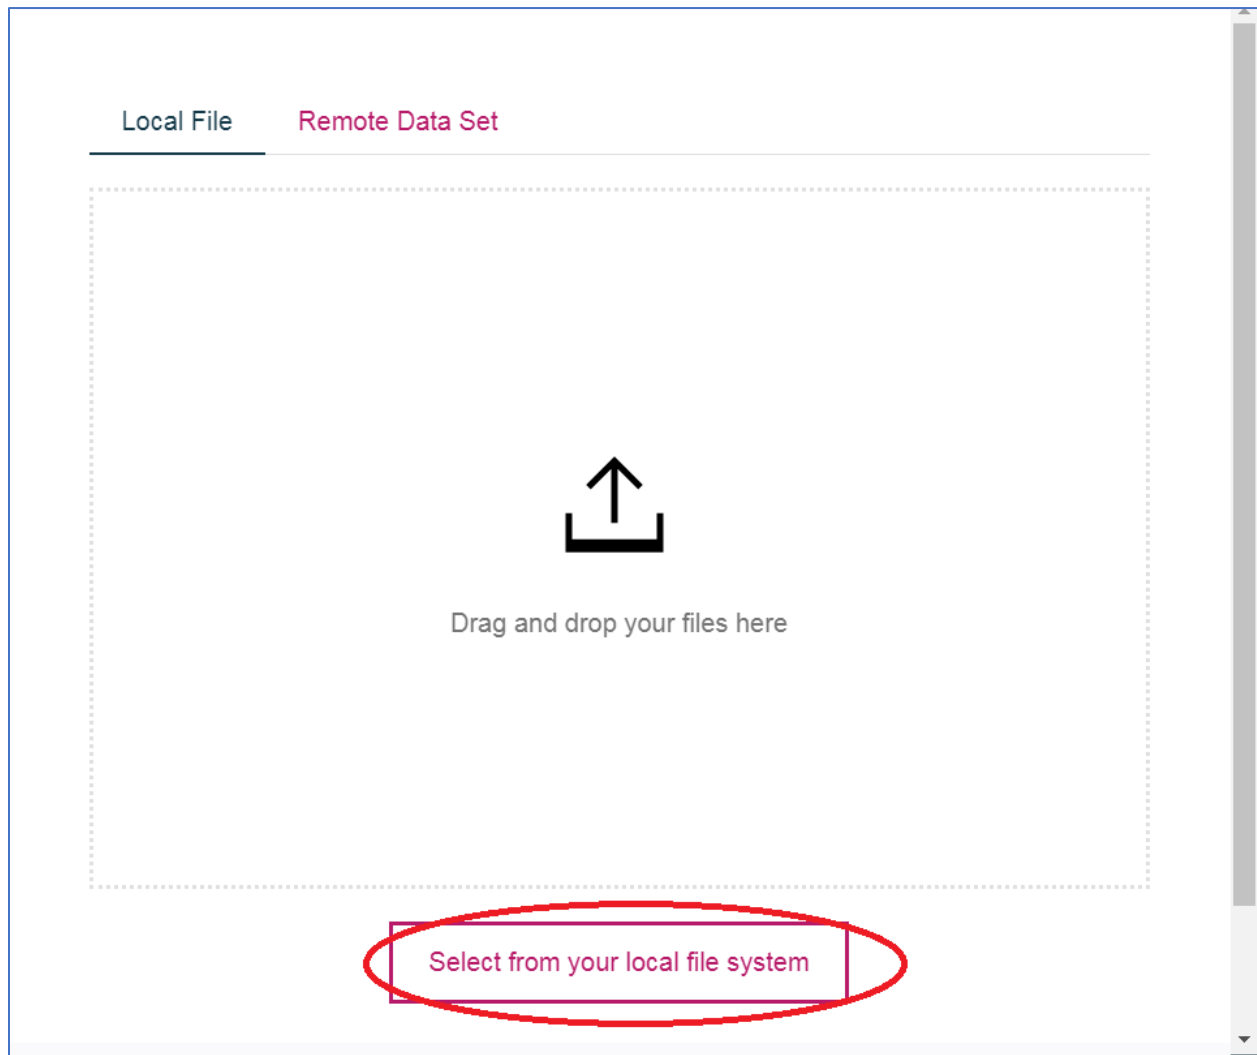
4. Click on **Data Sets**.



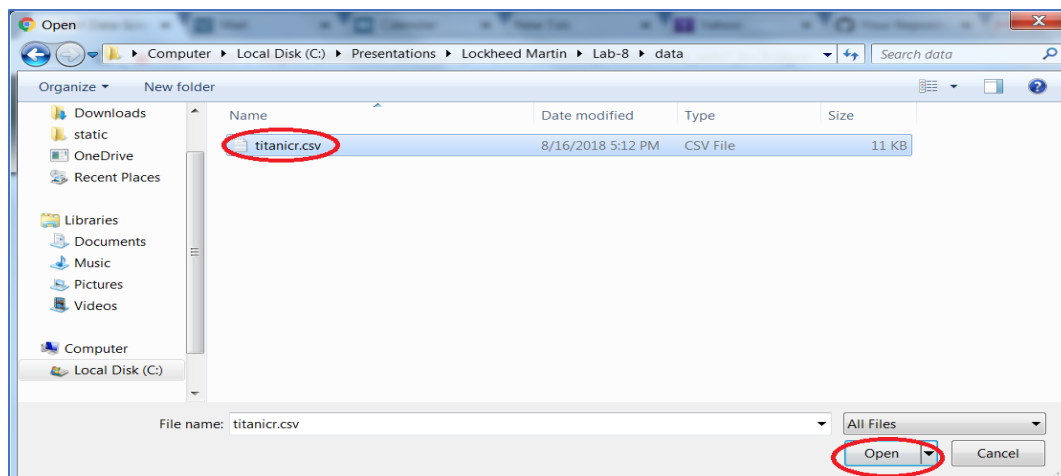
5. Click on **add data set**.



- Click on **Select from your local file system**.



- Navigate to the place where the **titanicr.csv** was downloaded. Select the **titanicr.csv** file and click **Open**.



8. The titanic.csv file is loaded into the project data sets.

Data Sets (33)					All	+	add data set
NAME	TYPE	SIZE	LAST MODIFIED	DATA SOURCE			
titanicr	CSV	10.42 KB	08-17-2018	Local File			
titanic	CSV	131.61 KB	08-16-2018	Local File			

Step 2: Profile the data to help determine missing values.

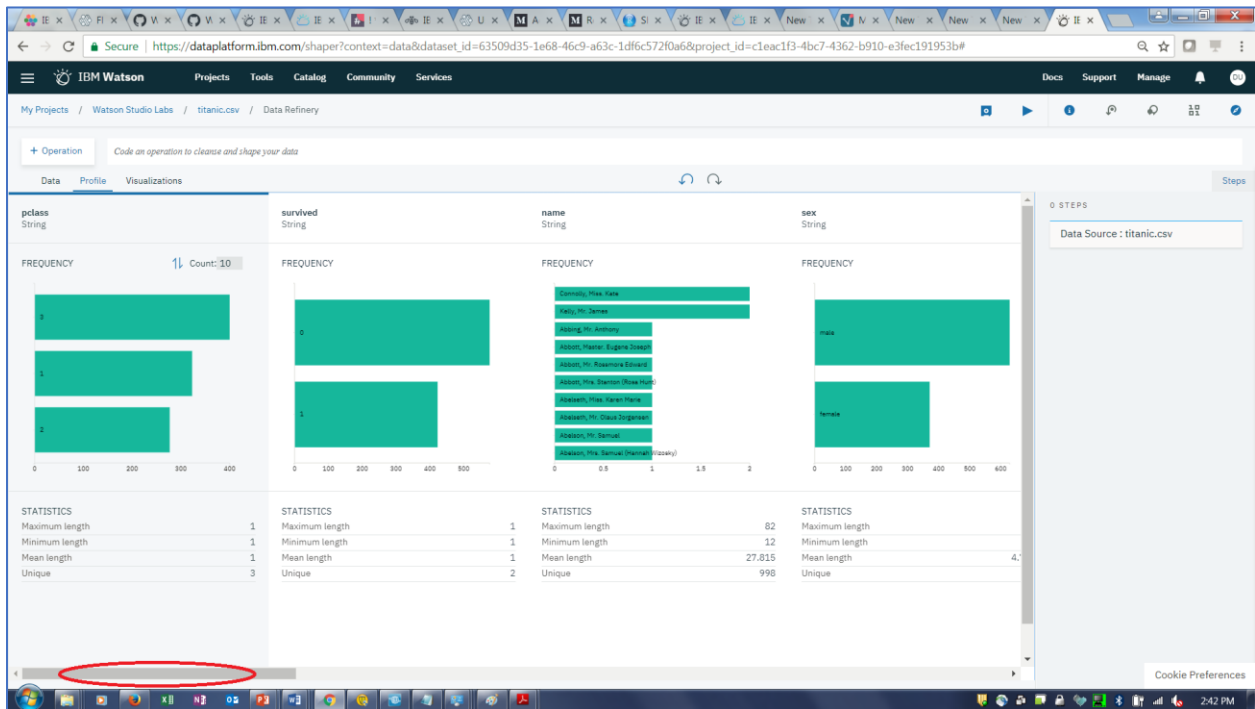
1. Add a Data Flow by clicking on the vertical ellipse and then click on **Refine**.

Data Sets (33)					All	+	add data set
NAME	TYPE	SIZE	LAST MODIFIED	DATA SOURCE			
titanicr	CSV	10.42 KB	08-17-2018	Local File			
titanic	CSV	131.61 KB	08-16-2018	Local File			
TelcoModelEval	CSV	35.8 KB	08-15-2018	Local File			
History_Transactions_v4	CSV	5.97 MB	08-14-2018	Local File			
Current_Transactions_v6	CSV	4.84 MB	08-14-2018	Local File			

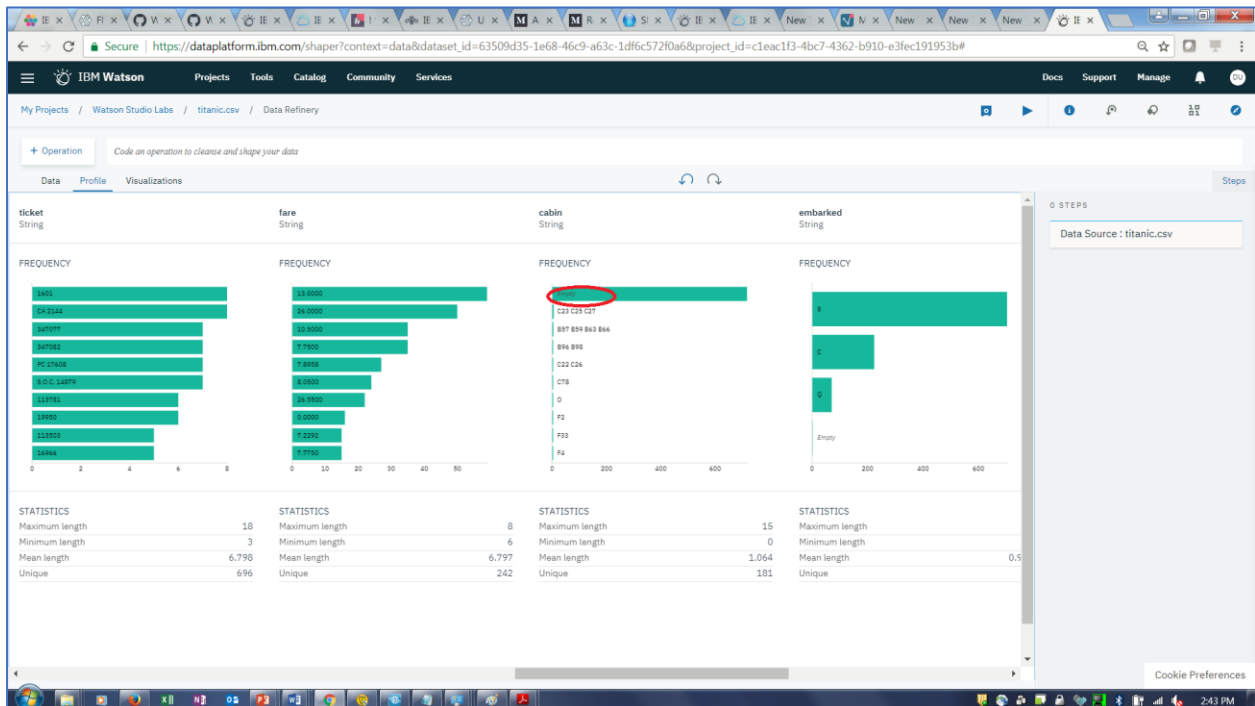
2. The Data Refinery panel will display the data set. Click on the **Profile** tab.

Data Refinery										
My Projects / Watson Studio Labs / titanic.csv / Data Refinery										
+ Operation Code an operation to cleanse and shape your data										
Data Profile Visualizations										
pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cal	Str
String	String	String	String	String	String	String	String	String	String	String
1	1	Allen, Miss. Elisabet...	female	29	0	0	24160	211.3375	BF	
2	1	Allison, Master. Hud...	male	0.9167	1	2	113781	151.5500	C2	
3	1	Allison, Miss. Helen ...	female	2	1	2	113781	151.5500	C2	
4	1	Allison, Mr. Hudson ...	male	30	1	2	113781	151.5500	C2	
5	1	Allison, Mrs. Hudso...	female	25	1	2	113781	151.5500	C2	
6	1	Anderson, Mr. Harry	male	48	0	0	19952	26.5500	E1	
7	1	Andrews, Miss. Korn...	female	63	1	0	13502	77.9583	D1	
8	1	Andrews, Mr. Thom...	male	39	0	0	112050	0.0000	A2	
9	1	Appleton, Mrs. Edw...	female	53	2	0	11769	51.4792	C1	
10	1	Artagaveytia, Mr. Ra...	male	71	0	0	PC 17609	49.5042	C2	
11	1	Astor, Col. John Jacob	male	47	1	0	PC 17757	227.5250	C6	
12	1	Astor, Mrs. John Jac...	female	18	1	0	PC 17757	227.5250	C6	
13	1	Aubart, Mme. LeontL...	female	24	0	0	PC 17477	69.3000	B5	
14	1	Barber, Miss. Ellen ...	female	26	0	0	19877	78.8500	A2	
15	1	Barkworth, Mr. Alge...	male	80	0	0	27042	30.0000	A2	
16	1	Baumann, Mr. John D	male		0	0	PC 17318	25.9250	C2	
17	1	Baxter, Mr. Quigg Ed...	male	24	0	1	PC 17558	247.5208	B1	
18	1	Baxter, Mrs. James (...)	female	50	0	1	PC 17558	247.5208	B1	
19	1	Bazzani, Miss. Albina	female	32	0	0	11813	76.2917	D1	
20	1	Beattie, Mr. Thomson	male	36	0	0	13050	75.2417	C6	

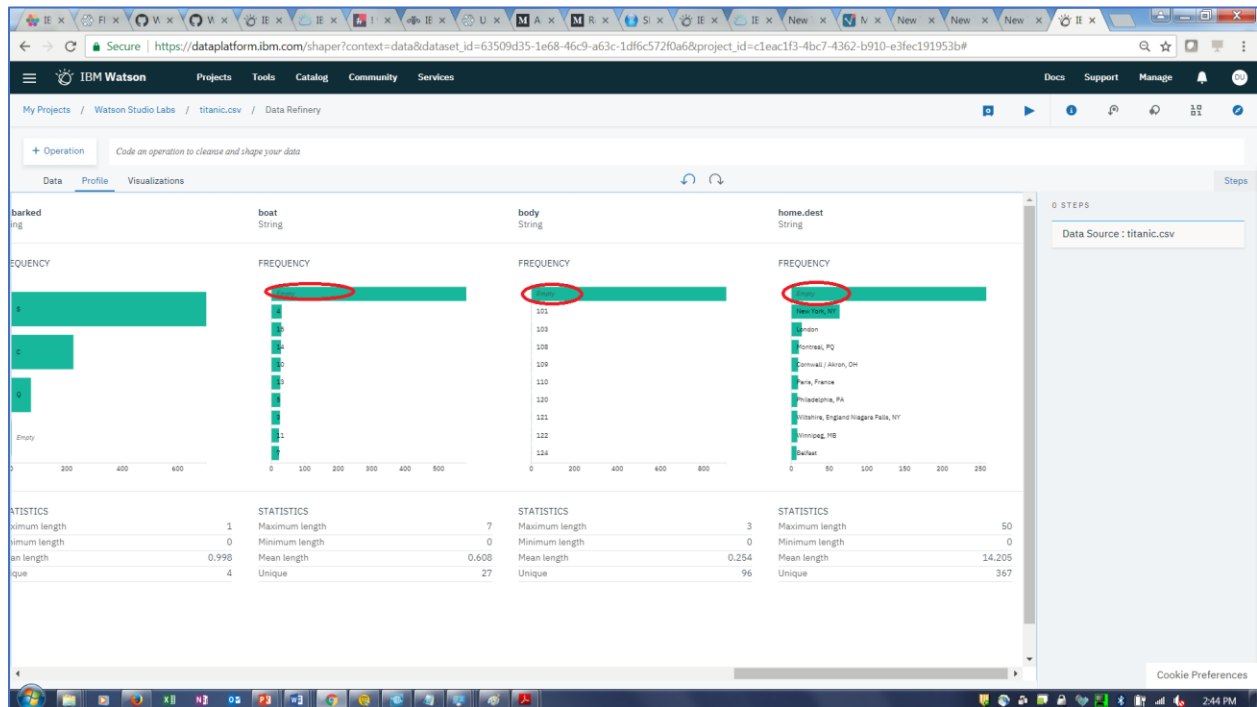
- The Profile panel displays the counts of the top 10 count values for each column. Note that you can change 10 to another number if desired. You can also switch to the bottom 10 counts for a column. Scroll to the right to view the cabin column.




- Note that the cabin column has many missing values and should be removed as part of the data preparation step.



5. In a similar fashion, scroll to the right to examine the boat, body, and home.dest columns. These also have many missing values and should be removed as part of the data preparation step.




6. Age and Embarked also have missing values. Embarked has only 1 missing value. Age has over 10 missing values, but we will keep that column in the analysis. As part of data preparation, we will remove the rows that contain the missing age and embarked values.
7. Click on the **Data** tab. We will add columns that contain more readable values for the survived and pclass columns. The column survived_value will contain a “Y” or “N”. The pclass_value column will contain “first”, “second”, or “third”. We will use the mutate (R dplyr function) and ifelse functions to do the conversion. Click on the **Code an operation to cleanse and shape your data**.




8. Type the following:
`mutate(survived_value=ifelse(survived==1, "Y", "N"))`

and then click Apply. If you scroll to the right you should see the new column “survived_value”.



- Type the following to create pclass_value,
mutate(pclass_value=ifelse(pclass==1,"first",ifelse(pclass==2,"second","third")))




- The result is shown below. Notice that the right panel will contain a running list of the transformations.

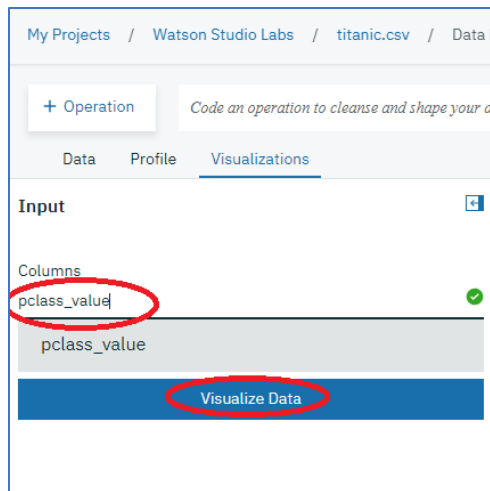
	ticket	fare	cabin	embarked	boat	body	home.dest	survived_value	pclass_value
	String	String	String	String	String	String	String	String	String
1	24160	211.3375	B5	S	2		St Louis, MO	Y	first
2	113781	151.5500	C22 C26	S	11		Montreal, PQ / Ches...	Y	first
3	113781	151.5500	C22 C26	S			Montreal, PQ / Ches...	N	first
4	113781	151.5500	C22 C26	S		135	Montreal, PQ / Ches...	N	first
5	113781	151.5500	C22 C26	S			Montreal, PQ / Ches...	N	first
6	19952	26.5500	E12	S	3		New York, NY	Y	first
7	13502	77.9583	D7	S	10		Hudson, NY	Y	first
8	112050	0.0000	A36	S			Belfast, NI	N	first
9	11769	51.4792	C101	S	0		Bayside, Queens, NY	Y	first
10	PC 17609	49.5042		C		22	Montevideo, Uruguay	N	first
11	PC 17757	227.5250	C62 C64	C		124	New York, NY	N	first
12	PC 17757	227.5250	C62 C64	C	4		New York, NY	Y	first

Step 3: Visualize the data to get a better understanding

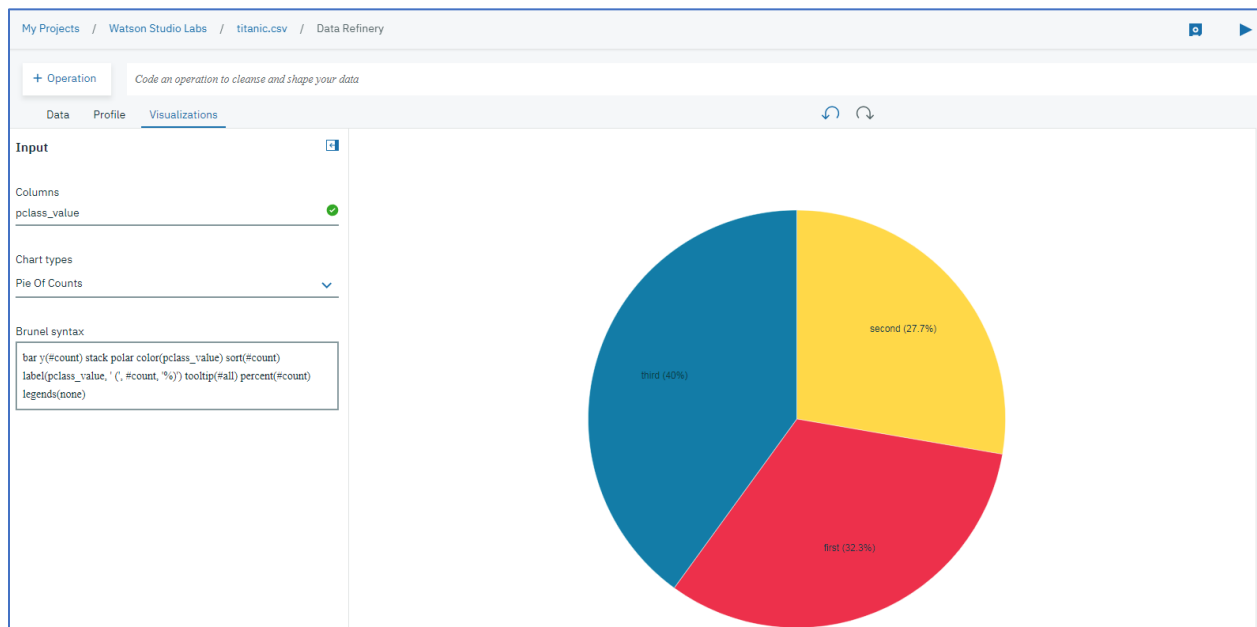
- Click on the **Visualizations** tab.



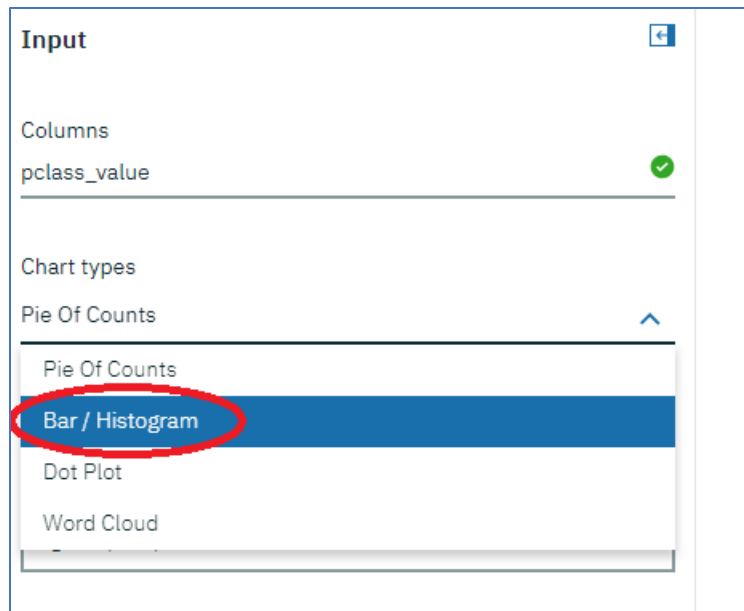
- Let's take a look at the breakdown of passengers by passenger class. We will use our new pclass_value field. Enter or select pclass_value and then click **Visualize Data**



3. The result is shown below.

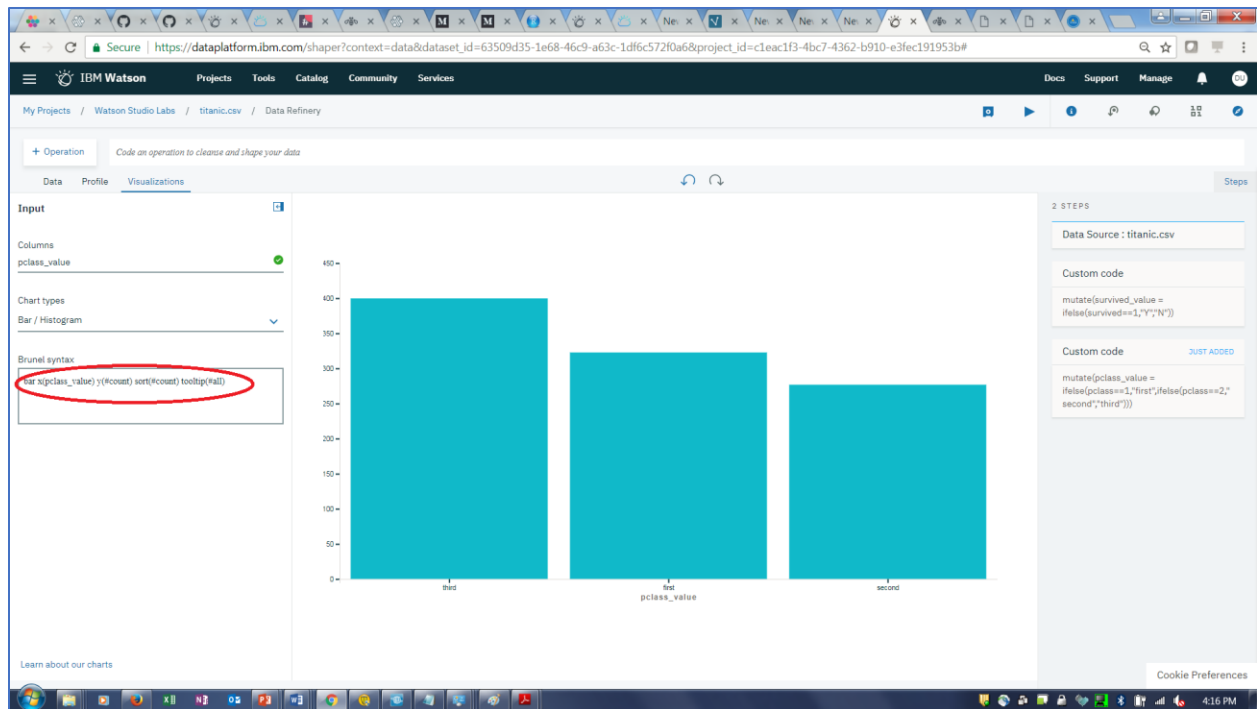


4. We can switch this to a bar chart, by switching the Chart type.

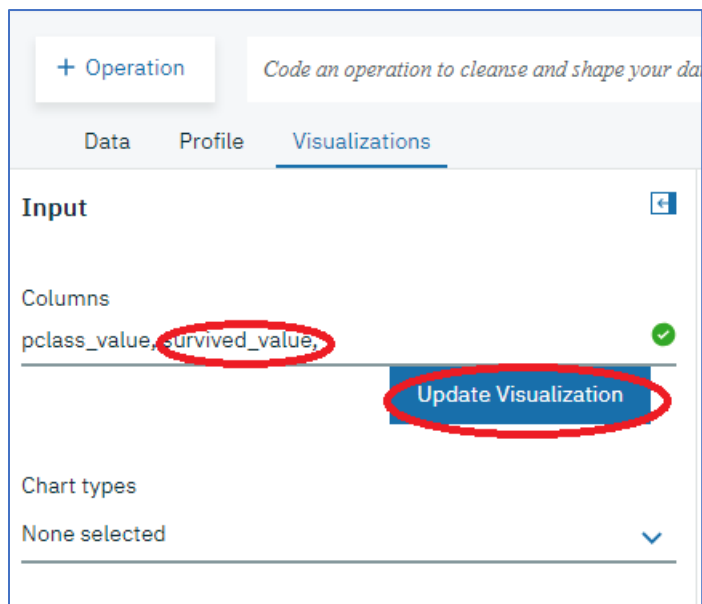


5. The result is shown below. Note the Brunel coding syntax. According to the Brunel github repo, *Brunel defines a highly succinct and novel language that defines interactive data visualizations based on tabular data. The language is well suited for both data scientists and more aggressive business users. The system interprets the language and produces visualizations using the user's choice of existing lower-level visualization technologies typically used by application engineers such as RAVE or D3. It can operate stand-alone and integrated into Jupyter (IPython) notebooks with further integrations as well as other low-level rendering support depending on the desires of the community.*


If you understand the syntax, you can make changes and update the visualization.



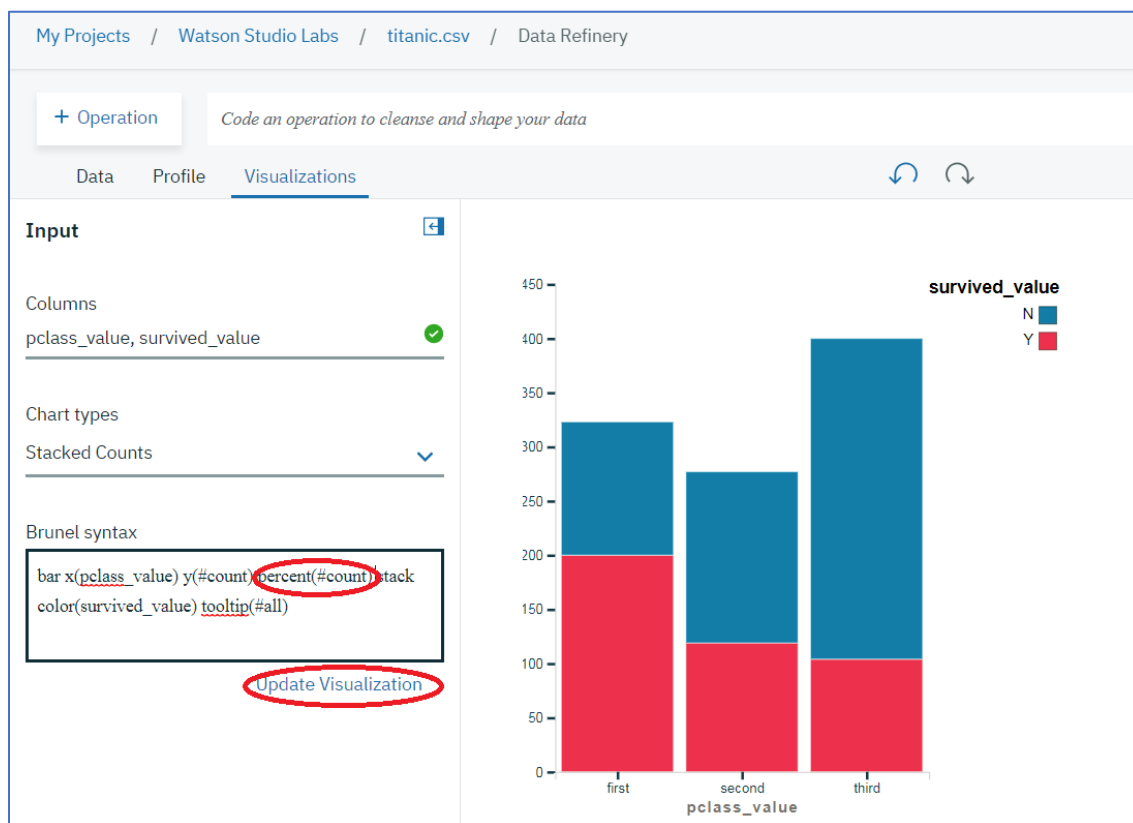
6. Let's examine the relationship between survival and the passenger class. We will add the `survived_value` and click **Update Visualization**.



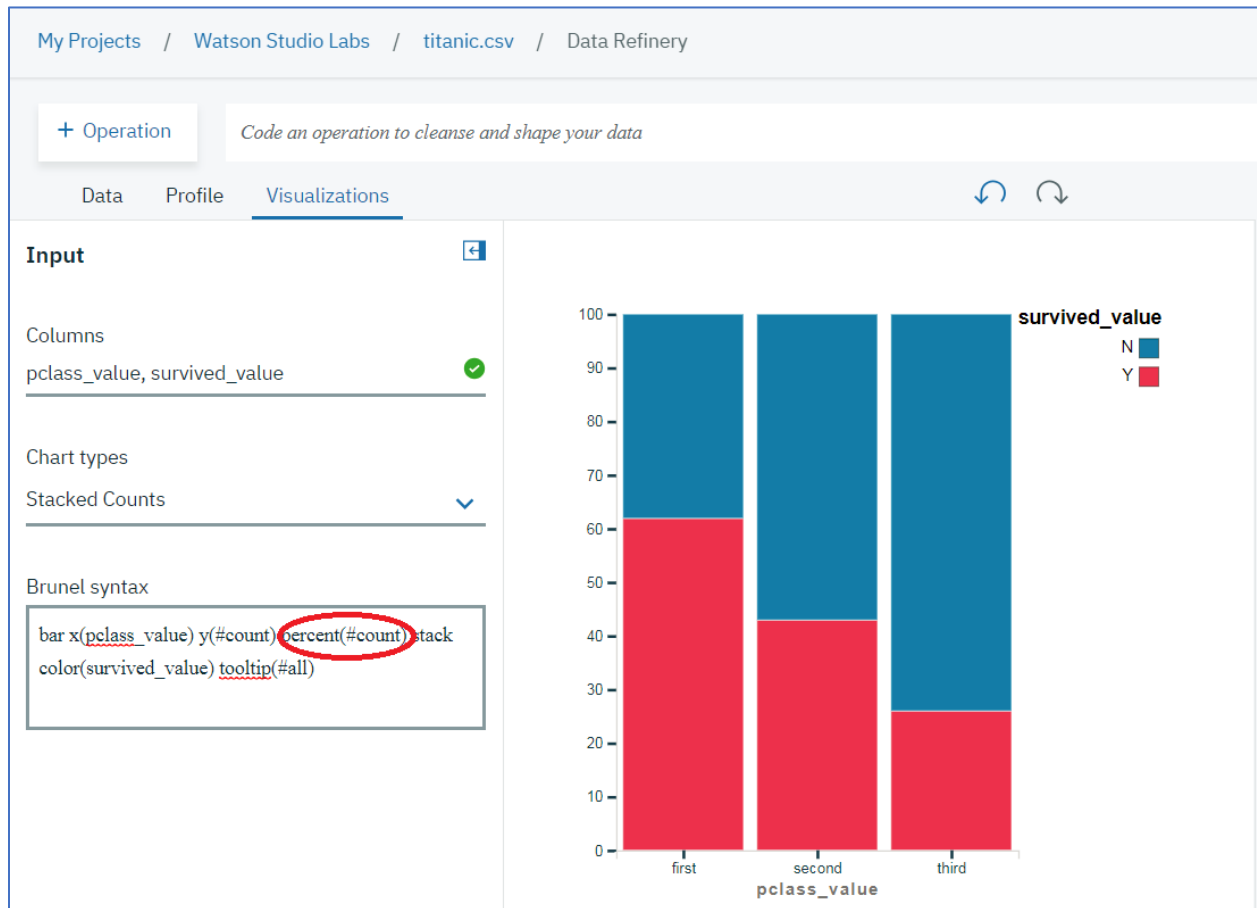
7. The result is shown below. We can see that survival probability for first class customers is significantly better.



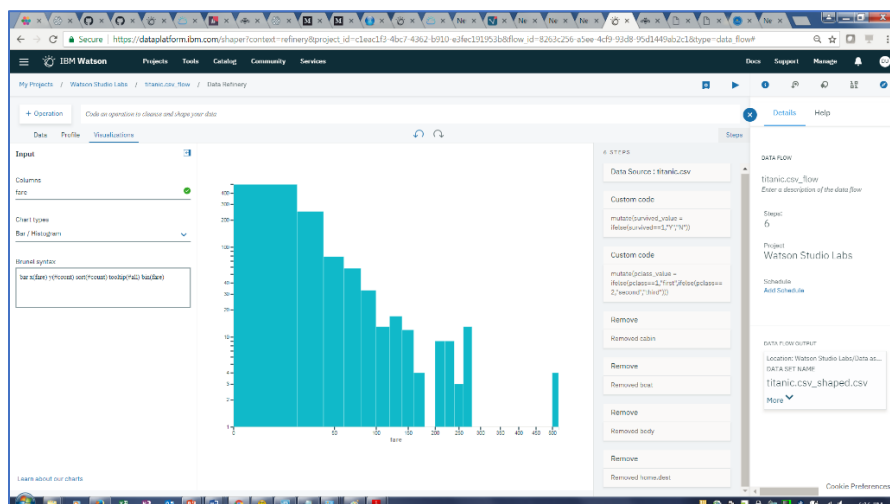
- If we want to normalize the results so that each column is shown as a percentage to allow comparisons more easily, then add **percent(#count)** to the Brunel syntax and click on **Update Visualization**.



9. The result is shown below. We can see that the percentage of survival is greatest for first class and lowest for third class.



10. Plot the fare values. The result is shown below. Note that it is highly skewed which affects the performance of some machine learning algorithms. One way to deal with this is to apply a logarithmic transformation. We will do that as part of data preparation.




Step 4: Prepare the data for modeling

Based on the data analysis, we need to do the following to prepare the data for modeling.

1. Remove columns cabin, boat, body, home.dest
2. Remove rows with missing values of age, and embarked.
3. Create a new column(log_fare) that is the logarithm of the fare column

We will also bin the age, and log_fare fields.

1. Return to the Data panel by clicking on the **Data** tab




2. Remove the cabin column by selecting on the vertical ellipse and then clicking on **Remove**.

The screenshot shows a table view in the Data Refinery interface. The table has three columns: 'cabin' (String), 'embarked' (String), and 'boat' (String). The 'cabin' column is selected, and a context menu is open over it. The menu options are: 'Remove' (highlighted with a red circle), 'Remove duplicates', 'Remove empty rows', 'Sort ascending', 'Sort descending', 'Substitute', 'CONVERT COLU... >', 'TEXT >', and 'View All'. The table data includes rows with cabin numbers like B5, C22 C26, E12, D7, A36, C101, C62 C64, and B35, along with their corresponding embarked status and boat numbers.

cabin String	embarked String	boat String
B5		2
C22 C26		11
C22 C26		
C22 C26		
C22 C26		
E12		3
D7		10
A36		
C101		D
C62 C64		
C62 C64	C	4
B35	C	9
	S	6

3. Remove the boat, body, and home.dest columns in a similar manner by selecting on the vertical ellipse adjacent to the column and clicking on **Remove**. Notice the STEPS panel on the right-hand side that provides a running list of the data operations.



4. For the age and embarked columns, click on the vertical ellipse adjacent to the columns, and click on **Remove empty rows**.

embarked	survived_value	pclass
String	String	String
S		first
S		first
S		first
S		first
S		first
S		first
S		first
S		first
S		first
C		first
C		first
C	Y	first
C	Y	first
S	Y	first

- Convert the fare column from a String to a Decimal by clicking on the vertical ellipse adjacent to the column, click on **Convert Column**, and then click on **Decimal**.

fare	embarked	survive	6 STEPS
String	String	String	
211.3375		Y	Data Sour
151.5500		Y	Custom co
151.5500		N	mutate(sur
151.5500		N	ifelse(survi
151.5500		N	
26.5500		Y	Custom co
77.9583		Y	mutate(pcl
0.0000		N	ifelse(pclas
51.4792			d",
49.5042			a
227.5250			d c
227.5250			d c
69.3000			a
78.8500			3
30.0000			d b

- Create a new column that is the log to the base 10 of the fare by clicking into the **Code an operation to cleanse and shape your data**, and entering

```
mutate(log_fare=log10(fare))
```

then click **Apply**.

+ Operation	mutate(log_fare=log10(fare))	Apply
-------------	------------------------------	-------

- Convert the age from String to Integer by clicking on the vertical ellipse adjacent to the age column, clicking on **Convert Column**, and clicking on **Integer**.


age	sibsp	parch	ticket
Integer	String	String	String
29		0	24160
0		2	11378
2		2	11378
30		2	11378
25		2	11378
48		0	19952
63		0	13502
39		0	11205
53			11769
71			PC 176
47	1		PC 177
18	1		PC 177
24	0		PC 174
26	0	0	19877

8. Bin the age column into the following bins by clicking into the **Code an operation to cleanse and shape your data**, and entering

```
mutate(age_bin=ifelse(age<6,0,ifelse(age<12,1,ifelse(age<18,2,ifelse(age<40,3,ifelse(age<65,4,ifelse(age<80,5,6)))))))
```

and then click **Apply**.


Bin	Age Range
0	0-5
1	6-11
2	12-17
3	18-39
4	40-64
5	65-79
6	Over 79



9. Bin the log_fare column, by clicking into the **Code an operation to cleanse and shape your data**, and entering

```
mutate(log_fare_bin=ifelse(log_fare<0,0,ifelse(log_fare>8,9,as.integer(log_fare)+1)))
```

and then clicking **Apply**




10. Now we will drop the age, fare, and log_fare columns as they are no longer needed for modeling purposes. Select the vertical ellipse adjacent to the column and click on **Remove** as shown below.

age	sibsp		
Integer	String		
29		Remove	
0		Remove duplicates	
2		Remove empty rows	
30		Sort ascending	
25		Sort descending	
48		Substitute	
63		CONVERT COLU... >	
39			
53		View All	

fare	embarked		
Decimal	String		
211.3375		Remove	
151.55		Remove duplicates	
151.55		Remove empty rows	
151.55		Sort ascending	
151.55		Sort descending	
26.55		Substitute	
77.9583		CONVERT COLU... >	
0			
51.4792		View All	
49.5042			
227.525	C		
227.525	C		

log_fare	age_bin
Decimal	Decimal
2.32497656566603	
2.18055594070364	
2.18055594070364	
2.18055594070364	
2.18055594070364	
1.42406452541749	
1.89186236009324	
-Inf	
1.71163178923691	
1.69464204659912	
2.35702912303943	4
2.35702912303943	3
1.84073323461181	3

11. Save the Data Flow by clicking on the Save Data Flow icon .



Save data flow


Save R-Script

embarked	survived_value
String	String
S	N


Step 5: Run the sequence of Data Flow operations on the entire data set.

When users are interacting with the Data Refinery tool, the operations are applied to a subset of the data set to facilitate faster response times. To run the data operations on the entire data set (titanic.csv), the user will generate an **R-Script** and run a **Job**.


1. Click on the  icon to **Save R-Script**.



2. Navigate to the project **Scripts** by clicking on the **DSX_Demos_XX** link and then clicking on **Scripts**.



3. Click on the vertical ellipse for **titanic.R** script, and click on **Create Job**.



4. Enter a **Name**, select **Script Run** for **Type**, select RStudio with **R 3.4.3 for Worker** for Worker, select the **Titanicr** script for **Source asset**, enter the input data set **datasets/titanic.csv** for **arg 1**, and enter the output data set **datasets/titanic_cleansed_out.csv** for **arg 2**, and then click **Create**.

Name *
TitanicDR 41

Description
Job description 300

Type *
Script run

Worker *
RStudio with R 3.4.3

Target host *
Local instance

Source asset *
/scripts/titanicR

Environment variables +
VARIABLE_1=value 1

Command line arguments +
datasets/titanic.csv
datasets/titanic_cleansed_out.csv

Cancel Create

5. Scroll down to the **Runs** section, and click **run now**.

Runs run now

ID	NAME	TARGET HOST	TRIGGERED BY	STARTED AT	DURATION (S)	RESULT
no runs found						

6. Enter a Name for the Run, and then click **Run**.

Run TitanicDR

Name *
Run1

Target host *
Local instance

Environment variables (+)
VARIABLE_1=value 1

Command line arguments (+)
datasets/titanic.csv
datasets/titanic_cleansed_out.csv

Cancel Run

7. Wait for the script to complete.

Run1 <div> Bernard Beekman 21 Aug 2018, 11:59 PM </div>			
No description available.			
RUN ID 1534910392-1001	TARGET HOST Local instance	DURATION 3 s	RESULT Success

8. When the script completes navigate to Data Sets by clicking on the **DSX_Local_XX** project link and then click on **Data Sets**. Notice that the **titanic_cleansed_out.csv** data set has been created.

All

Notebooks

RStudio

Models

SPSS Modeler Flows

Scripts

Data Sets

Other Files

Data Flows

Published Assets

Data Sets

(15)

All

+

add data set

NAME	TYPE	SIZE	LAST MODIFIED	DATA SOURCE
<div><div></div><div>titanic_cleansed_out</div></div>	CSV	78.49 KB	08-22-2018	Local File
<div><div></div><div>titanic</div></div>	CSV	105.73 KB	08-22-2018	Local File
<div><div></div><div>titanicr</div></div>	CSV	10.42 KB	08-22-2018	Local File

- Click on the **titanic_cleansed_out.csv** to preview the data set. We can see that the new fields defined in the Data Refinery flow have been created. Click on **Close** to remove the Preview panel.

Preview - titanic_cleansed_out.csv

id	ticket	embarked	survived_value	pclass_value	age_bin	log_fare_bin
	24160	S	Y	first	3	3
	113781	S	Y	first	0	3
	113781	S	N	first	0	3
	113781	S	N	first	3	3

Close